

Multithreading with modern C++

Jui-Hung Hung

Memory model of a process

- Process is a running program
- Every process has its virtual memory allocation

A Computer Process

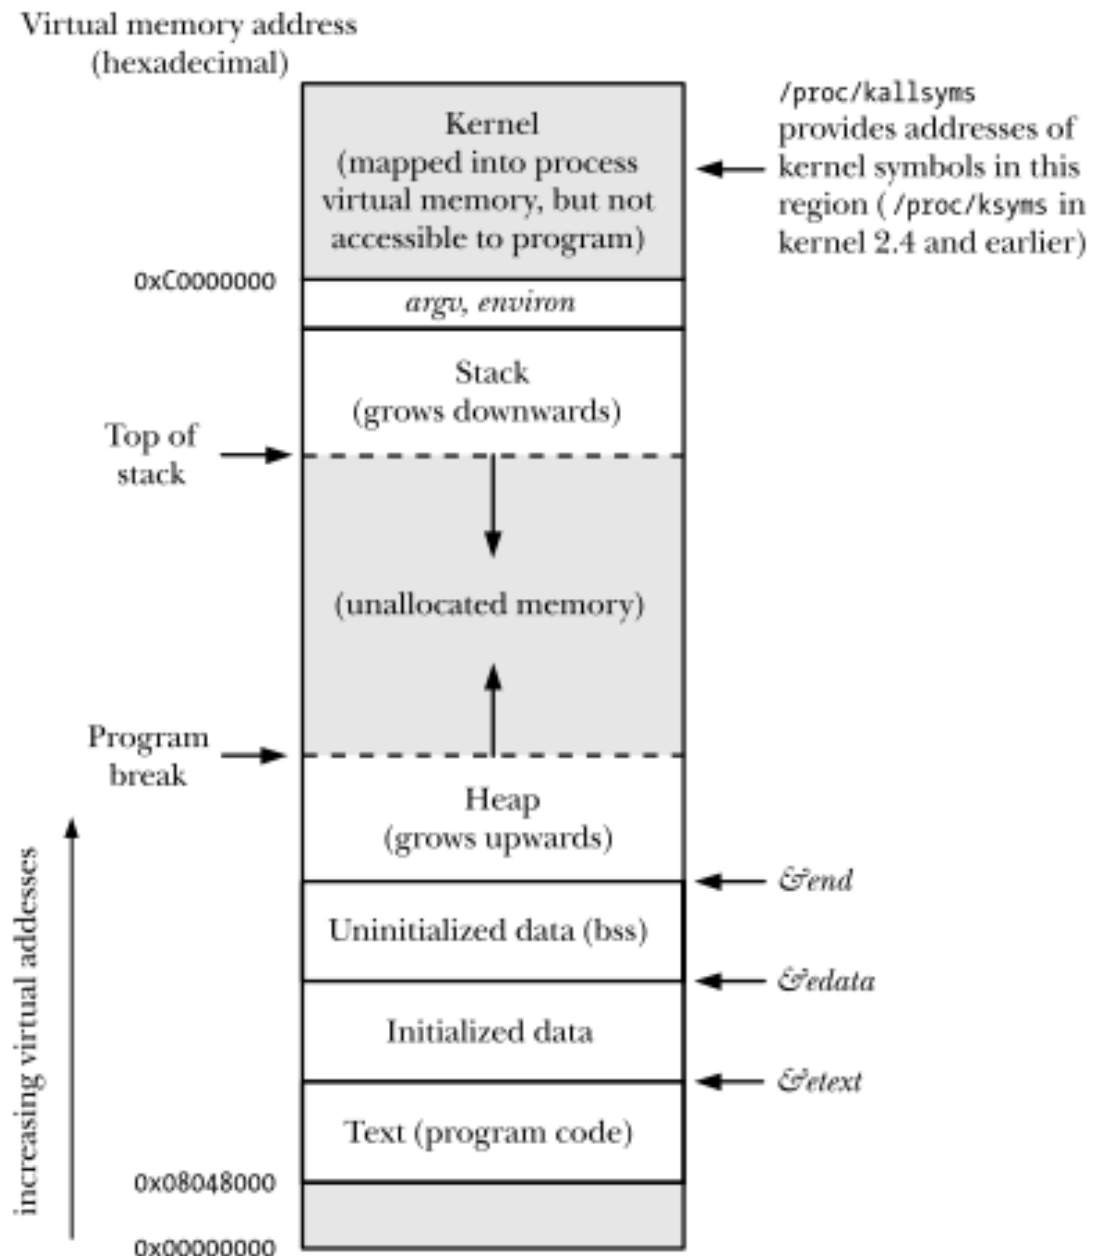
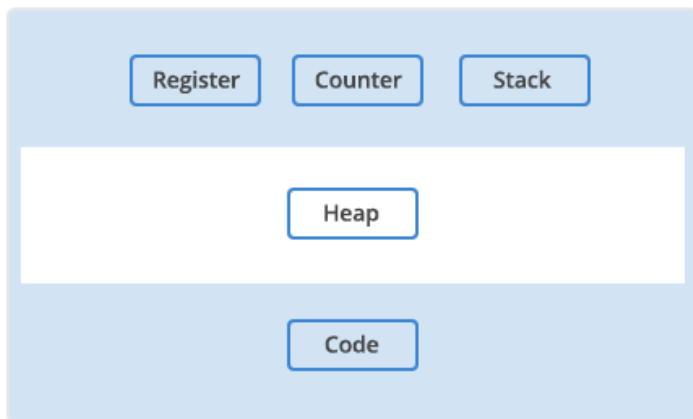
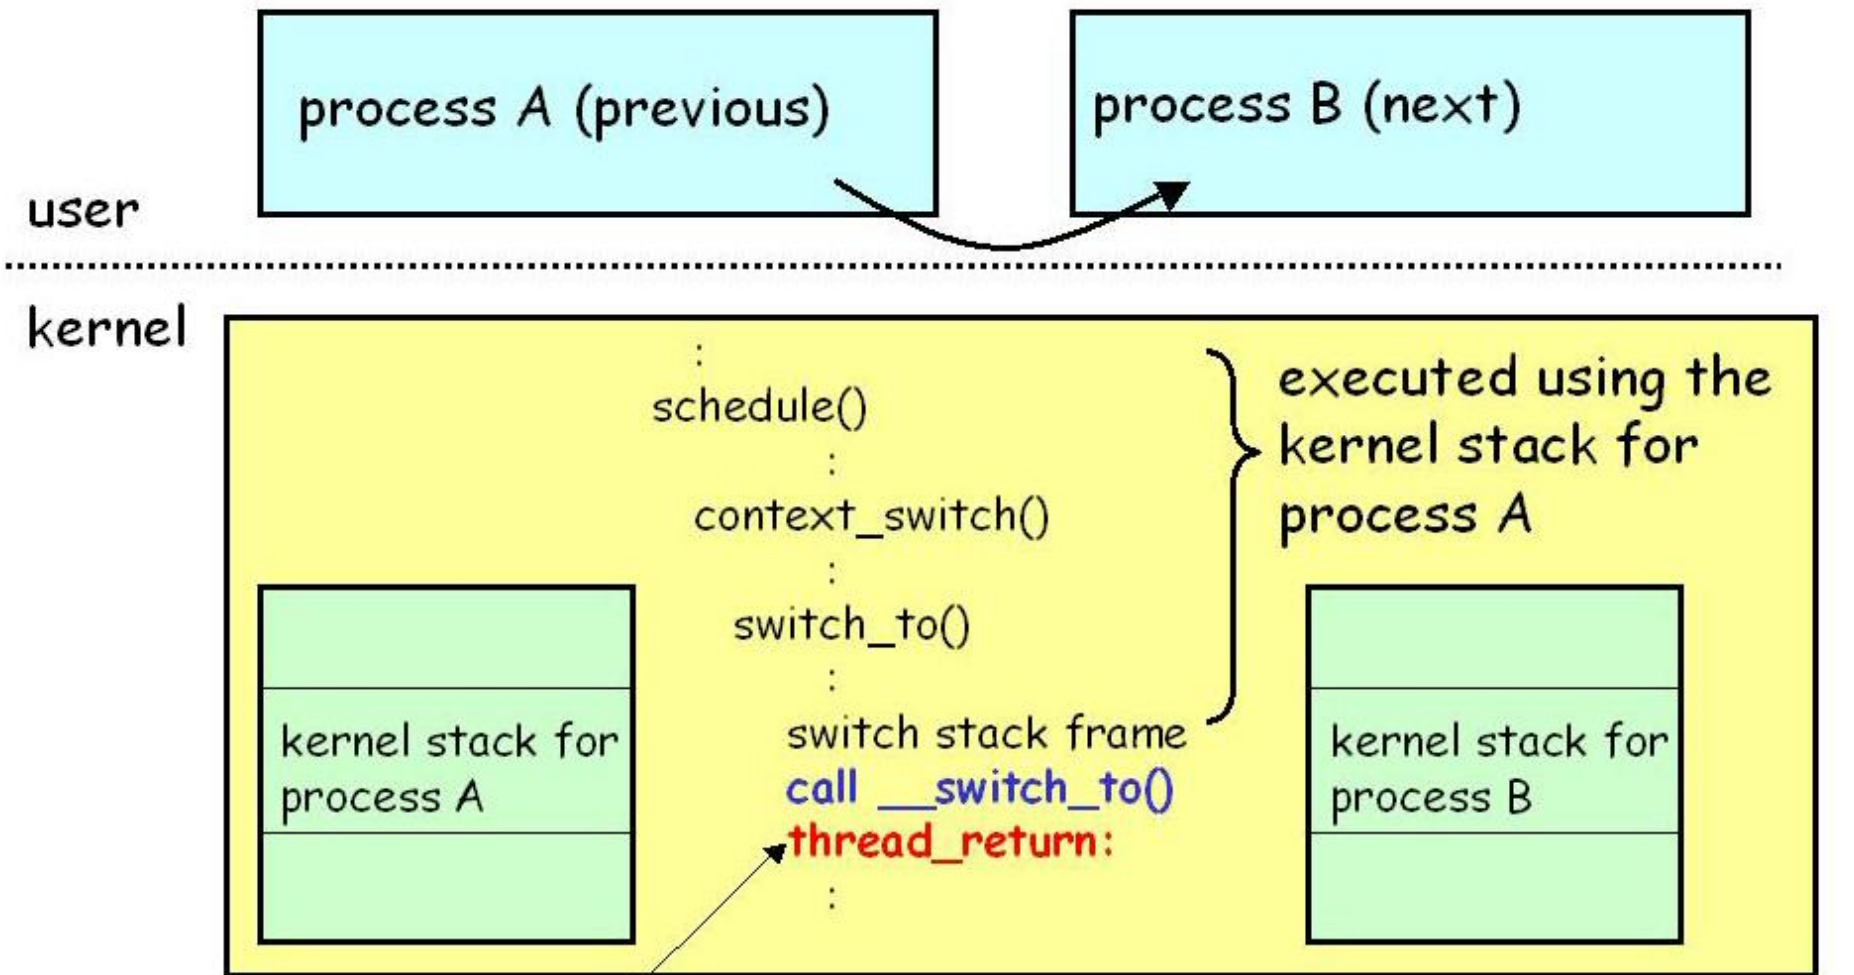


Figure 6-1: Typical memory layout of a process on Linux/x86-32

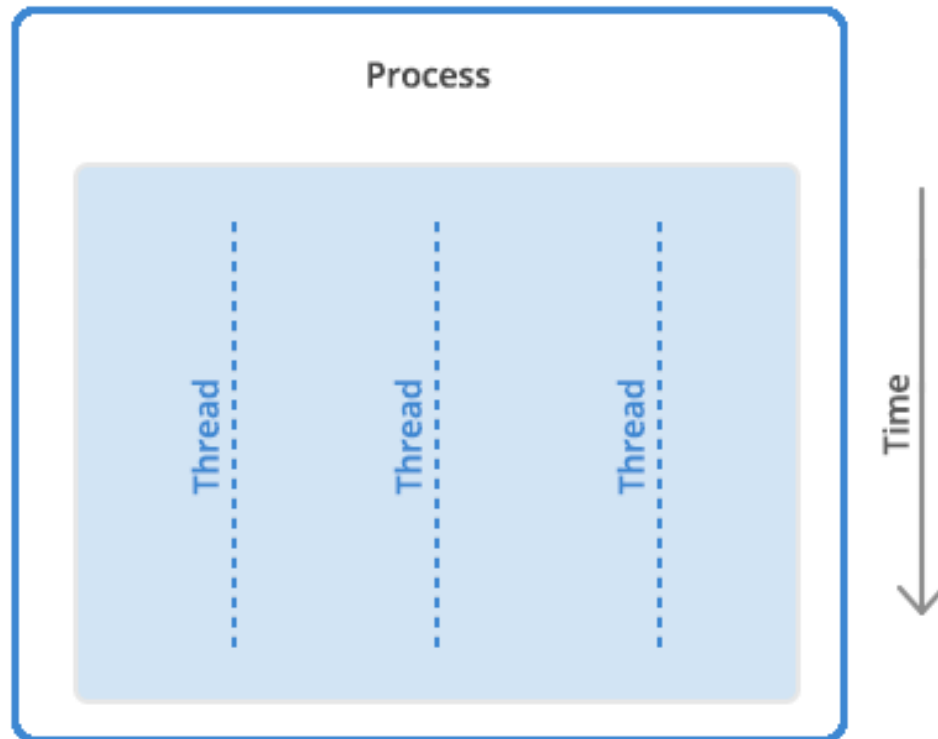
Context Switch of Linux(x86_64)



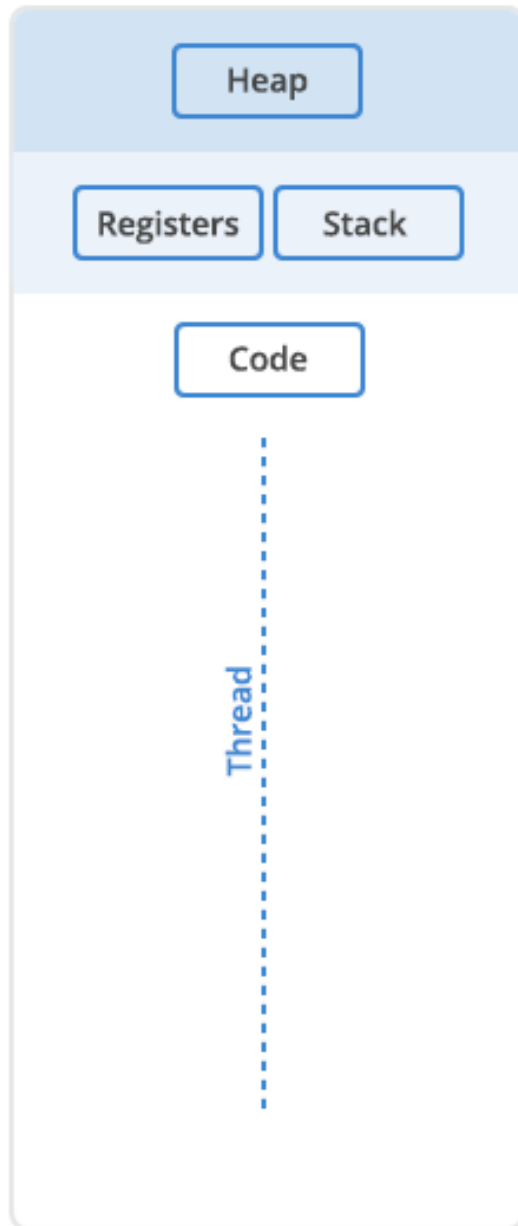
Stored as the return address by '`call __switch_to()`' into the kernel stack for **process B**. Thus, on return of `__switch_to()`, process B begins to run from `thread_return`.

Thread (lightweight process)

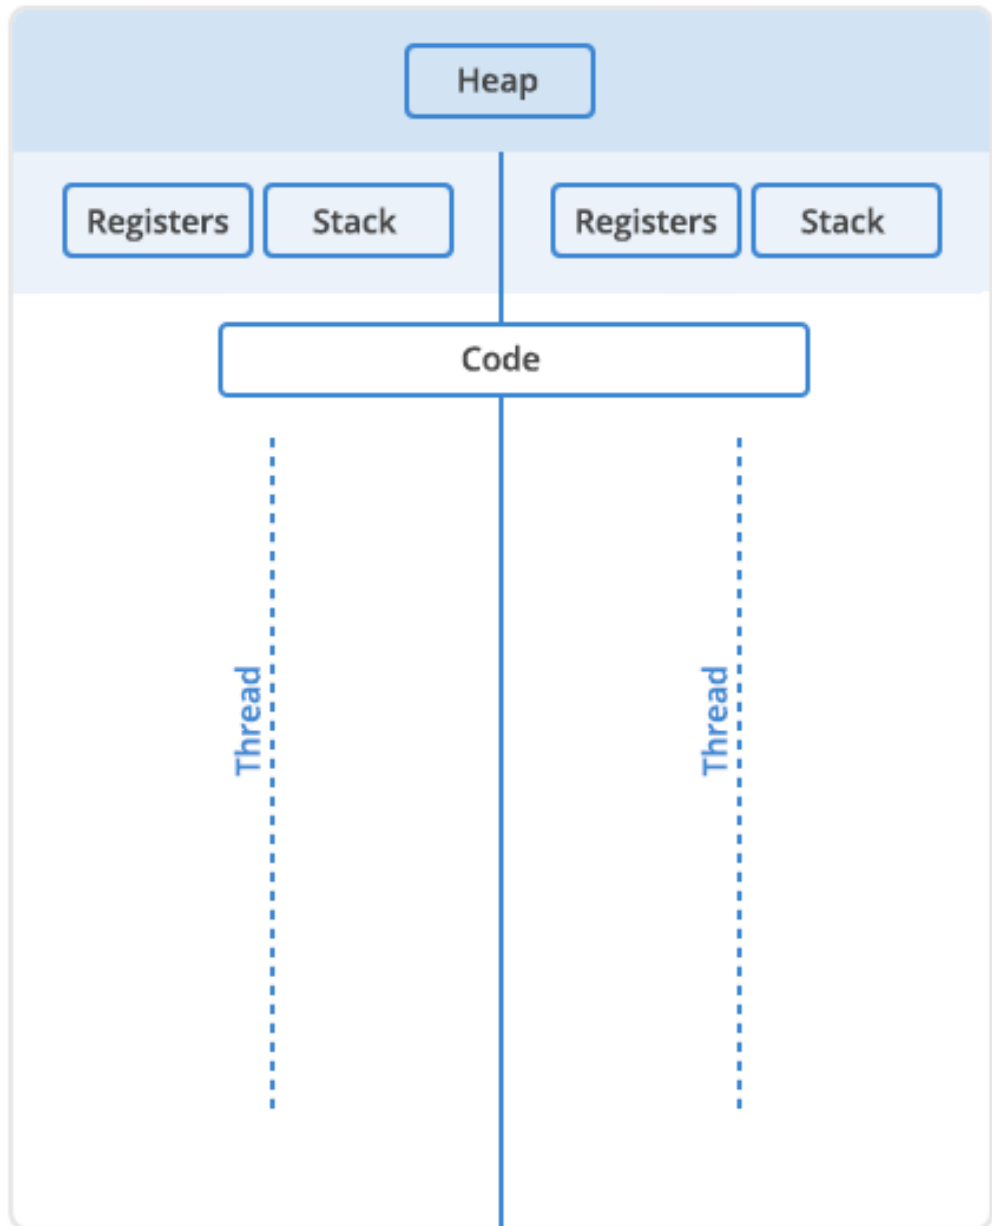
- A thread is the unit of execution within a process
- spreading a single process across multiple threads



Single Thread



Multi Threaded



Life cycle

1. The program starts out as a text file of programming code,
2. The program is compiled or interpreted into binary form,
3. The program is loaded into memory,
4. The program becomes one or more running processes.
5. Processes are typically independent of each other,
6. While threads exist as the subset of a process.
 1. Threads can communicate with each other more easily than processes can,
 2. But threads are more vulnerable to problems caused by other threads in the same process.

PROCESS

THREAD

Processes are heavyweight operations

Threads are lighter weight operations

Each process has its own memory space

Threads use the memory of the process they belong to

Inter-process communication is slow as processes have different memory addresses

Inter-thread communication can be faster than inter-process communication because threads of the same process share memory with the process they belong to

Context switching between processes is more expensive

Context switching between threads of the same process is less expensive

Processes don't share memory with other processes

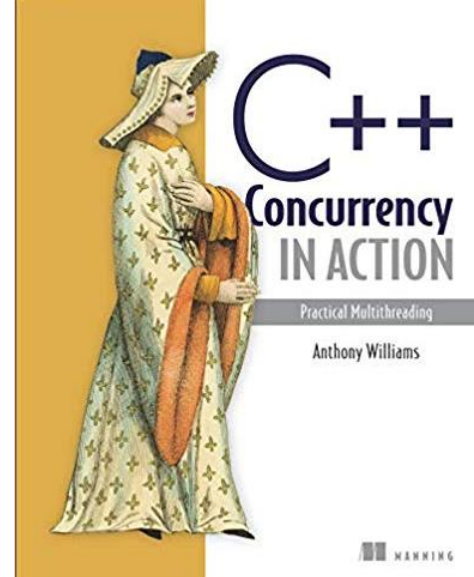
Threads share memory with other threads of the same process

Hello concurrent world!

```
#include <iostream>
#include <thread>          ← 1

void hello()              ← 2
{
    std::cout<<"Hello Concurrent World\n";
}

int main()
{
    std::thread t(hello);  ← 3
    t.join();              ← 4
}
```



Ways of invoking a thread

```
void do_some_work();  
std::thread my_thread(do_some_work);
```

```
std::thread my_thread([](  
    do_something();  
    do_something_else();  
));
```

```
class X  
{  
public:  
    void do_lengthy_work();  
};  
  
X my_x;  
std::thread t(&X::do_lengthy_work, &my_x);
```

```
class background_task  
{  
public:  
    void operator()() const  
    {  
        do_something();  
        do_something_else();  
    }  
};  
  
background_task f;  
std::thread my_thread(f);
```

Std::bind

```
#include <random>
#include <iostream>
#include <memory>
#include <functional>

void f(int n1, int n2, int n3, const int& n4, int n5)
{
    std::cout << n1 << ' ' << n2 << ' ' << n3 << ' ' << n4 << ' ' << n5 << '\n';
}

int g(int n1)
{
    return n1;
}

struct Foo {
    void print_sum(int n1, int n2)
    {
        std::cout << n1+n2 << '\n';
    }
    int data = 10;
};
```

The arguments to bind are copied or moved, and are never passed by reference unless wrapped in [std::ref](#) or [std::cref](#).

```

int main()
{
    using namespace std::placeholders; // for _1, _2, _3...

    // demonstrates argument reordering and pass-by-reference
    int n = 7;
    // (_1 and _2 are from std::placeholders, and represent future
    // arguments that will be passed to f1)
    auto f1 = std::bind(f, _2, 42, _1, std::cref(n), n);
    n = 10;
    f1(1, 2, 1001); // 1 is bound by _1, 2 is bound by _2, 1001 is unused
                  // makes a call to f(2, 42, 1, n, 7)

    // nested bind subexpressions share the placeholders
    auto f2 = std::bind(f, _3, std::bind(g, _3), _3, 4, 5);
    f2(10, 11, 12); // makes a call to f(12, g(12), 12, 4, 5);

    // common use case: binding a RNG with a distribution
    std::default_random_engine e;
    std::uniform_int_distribution<> d(0, 10);
    auto rnd = std::bind(d, e); // a copy of e is stored in rnd
    for(int n=0; n<10; ++n)
        std::cout << rnd() << ' ';
    std::cout << '\n';

    // bind to a pointer to member function
    Foo foo;
    auto f3 = std::bind(&Foo::print_sum, &foo, 95, _1);
    f3(5);

    // bind to a pointer to data member
    auto f4 = std::bind(&Foo::data, _1);
    std::cout << f4(foo) << '\n';

    // smart pointers can be used to call members of the referenced objects, too
    std::cout << f4(std::make_shared<Foo>(foo)) << '\n'
              << f4(std::make_unique<Foo>(foo)) << '\n';
}

```

Join or detach a thread

- Join
 - Wait for the thread to complete
 - cleans up any storage associated with the thread
- Detach
 - leaves the thread to run in the background, with no direct means of communicating with it.
 - It's no longer possible to wait for that thread to complete
 - The thread can no longer be joined.

Thread guard (auto joining at destruction)

```
class thread_guard
{
    std::thread& t;
public:
    explicit thread_guard(std::thread& t_):
        t(t_)
    {}
    ~thread_guard()
    {
        if(t.joinable())           ← 1
        {
            t.join();              ← 2
        }
    }
    thread_guard(thread_guard const&)=delete;           ← 3
    thread_guard& operator=(thread_guard const&)=delete;
};

struct func;

void f()
{
    int some_local_state=0;
    func my_func(some_local_state);
    std::thread t(my_func);
    thread_guard g(t);

    do_something_in_current_thread();
}
```

See definition
in listing 2.1

- local objects are destroyed in reverse order of construction.
- thread_guard object g is destroyed first, and the thread is joined with in the destructor

Passing parameter into a thread can be tricky

```
struct func
{
    int& i;

    func(int& i_):i(i_){}

    void operator() ()
    {
        for(unsigned j=0;j<1000000;++j)
        {
            do_something(i);
        }
    }
};

void oops()
{
    int some_local_state=0;
    func my_func(some_local_state);
    std::thread my_thread(my_func);
    my_thread.detach();
}
```

1 Potential access to dangling reference

```
void f(int i,std::string const& s);

void not_oops(int some_param)
{
    char buffer[1024];
    sprintf(buffer,"%i",some_param);
    std::thread t(f,3,std::string(buffer));
    t.detach();
}
```

by default the arguments are *copied* into internal storage

Use `std::ref` or `cref` to pass by reference

```
void update_data_for_widget(widget_id w, widget_data& data);    ← 1  
  
void oops_again(widget_id w)  
{  
    widget_data data;  
    std::thread t(update_data_for_widget, w, data);              ← 2  
    display_status();  
    t.join();  
    process_widget_data(data);    ← 3  
}
```

```
    std::thread t(update_data_for_widget, w, std::ref(data));
```

Thread is only movable not copyable

```
class scoped_thread
{
    std::thread t;
public:
    explicit scoped_thread(std::thread t_):           ← 1
        t(std::move(t_))
    {
        if(!t.joinable())                            ← 2
            throw std::logic_error("No thread");
    }
    ~scoped_thread()
    {
        t.join();                                     ← 3
    }
    scoped_thread(scoped_thread const&)=delete;
    scoped_thread& operator=(scoped_thread const&)=delete;
};

struct func;
void f()
{
    int some_local_state;
    scoped_thread t(std::thread(func(some_local_state))); ← 4
    do_something_in_current_thread();
}                                                         ← 5
```

See
listing 2.1

Hardware concurrency

```
unsigned long const hardware_threads=  
    std::thread::hardware_concurrency();  
  
std::thread::id master_thread;  
void some_core_part_of_algorithm()  
{  
    if(std::this_thread::get_id()==master_thread)  
    {  
        do_master_thread_work();  
    }  
    do_common_work();  
}
```

Sharing data between threads

- *Care-free data/function call*
 - *read-only, there's no problem, because the data read by one thread is unaffected by whether or not another thread is reading the same data.*
 - *invariants*—statements that are always true about a particular data structure, such as “this variable contains the number of items in the list”
- Data race
 - the outcome depends on the relative ordering of execution of operations on two or more threads
 - cause the dreaded *undefined behavior*
 - can be benign or *problematic*

Two solutions for data race

- Finest granularity
 - modify the design of your data structure and its invariants so that modifications are done as a series of **indivisible** changes, each of which preserves the invariants. This is generally referred to as *lock-free programming*
- Protect critical sections
 - wrap your data structure with a **protection** mechanism, to ensure that only the thread actually performing a modification can see the intermediate states where the invariants are broken

Data race

```
#include <list>
#include <mutex>
#include <algorithm>

std::list<int> some_list;

void add_to_list(int new_value)
{
    some_list.push_back(new_value);
}

bool list_contains(int value_to_find)
{
    return std::find(some_list.begin(), some_list.end(), value_to_find)
        != some_list.end();
}
```

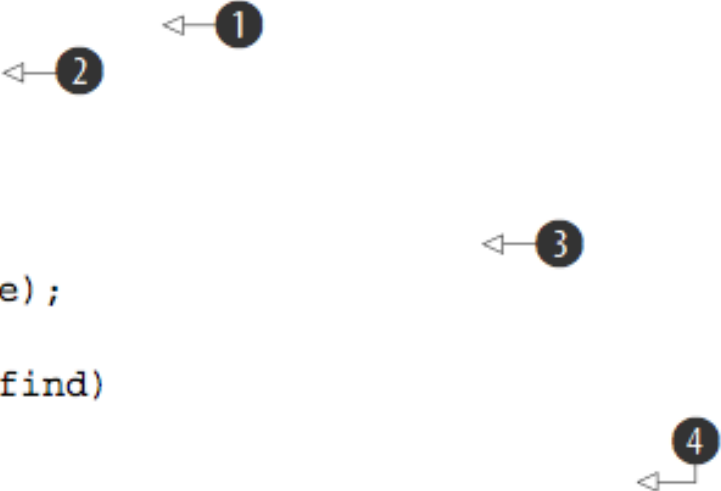


Diagram illustrating a data race scenario in the provided C++ code:

- 1. Declaration of `some_list` as a `std::list<int>`.
- 2. Access to `some_list` within the `add_to_list` function.
- 3. Modification of `some_list` via `push_back` within the `add_to_list` function.
- 4. Access to `some_list` within the `list_contains` function.

The code shows two functions, `add_to_list` and `list_contains`, both accessing the shared variable `some_list` without mutual exclusion (mutex), leading to a data race.

Mutex and unique_lock

```
#include <list>
#include <mutex>
#include <algorithm>
```

```
std::list<int> some_list;
std::mutex some_mutex;
```

```
void add_to_list(int new_value)
```

```
{
    std::lock_guard<std::mutex> guard(some_mutex);
    some_list.push_back(new_value);
}
```

```
bool list_contains(int value_to_find)
```

```
{
    std::lock_guard<std::mutex> guard(some_mutex);
    return std::find(some_list.begin(), some_list.end(), value_to_find)
        != some_list.end();
}
```

Critical section; access shared data / data race
(define by the life cycle of a lock)

Mutex (door to the critical section)

Lock (own the mutex)



3

4

AVOID NESTED LOCKS

don't wait for another thread if there's a chance it's waiting for you

Jailbreak

```
class some_data
{
    int a;
    std::string b;
public:
    void do_something();
};
```

```
class data_wrapper
{
private:
    some_data data;
    std::mutex m;
public:
    template<typename Function>
    void process_data(Function func)
    {
        std::lock_guard<std::mutex> l(m);
        func(data);
    }
};
```

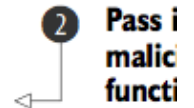
```
some_data* unprotected;

void malicious_function(some_data& protected_data)
{
    unprotected=&protected_data;
}

data_wrapper x;

void foo()
{
    x.process_data(malicious_function);
    unprotected->do_something();
}
```

2 Pass i
malic
functi



Thread-safe data structure

Thread A	Thread B
<pre>if(!s.empty()) int const value=s.top(); s.pop(); do_something(value);</pre>	<pre>if(!s.empty()) int const value=s.top(); s.pop(); do_something(value);</pre>

A thread-safe stack

```
#include <exception>
#include <memory>

struct empty_stack: std::exception
{
    const char* what() const throw();
};

template<typename T>
class threadsafe_stack
{
public:
    threadsafe_stack();
    threadsafe_stack(const threadsafe_stack&);
    threadsafe_stack& operator=(const threadsafe_stack&) = delete;

    void push(T new_value);
    std::shared_ptr<T> pop();
    void pop(T& value);
    bool empty() const;
};
```

← For `std::shared_ptr<>`

1
Assignment
operator is
deleted

←


```

private:
    std::stack<T> data;
    mutable std::mutex m;
public:
    threadsafe_stack(){}
    threadsafe_stack(const threadsafe_stack& other)
    {
        std::lock_guard<std::mutex> lock(other.m);
        data=other.data;
    }
    threadsafe_stack& operator=(const threadsafe_stack&) = delete;

    void push(T new_value)
    {
        std::lock_guard<std::mutex> lock(m);
        data.push(new_value);
    }
    std::shared_ptr<T> pop()
    {
        std::lock_guard<std::mutex> lock(m);
        if(data.empty()) throw empty_stack();
        std::shared_ptr<T> const res(std::make_shared<T>(data.top()));
        data.pop();
        return res;
    }
    void pop(T& value)
    {
        std::lock_guard<std::mutex> lock(m);
        if(data.empty()) throw empty_stack();
        value=data.top();
        data.pop();
    }
    bool empty() const
    {
        std::lock_guard<std::mutex> lock(m);
        return data.empty();
    }

```

1 Copy performed in constructor body

Check for empty before trying to pop value

Allocate return value before modifying stack

Deadlock

- This is almost the opposite of a race condition: rather than two threads racing to be first, each one is waiting for the other, so neither makes any progress
- The common advice for avoiding deadlock is to always lock the two mutexes in the same order: if you always lock mutex A before mutex B, then you'll never deadlock.

Avoid deadlock

std::defer_lock, std::try_to_lock, std::adopt_lock

Defined in header <mutex>

constexpr std::defer_lock_t defer_lock {};	(since C++11) (until C++17)
inline constexpr std::defer_lock_t defer_lock {};	(since C++17)
constexpr std::try_to_lock_t try_to_lock {};	(since C++11) (until C++17)
inline constexpr std::try_to_lock_t try_to_lock {};	(since C++17)
constexpr std::adopt_lock_t adopt_lock {};	(since C++11) (until C++17)
inline constexpr std::adopt_lock_t adopt_lock {};	(since C++17)

```
class some_big_object;
void swap(some_big_object& lhs, some_big_object& rhs);

class X
{
private:
    some_big_object some_detail;
    std::mutex m;
public:
    X(some_big_object const& sd):some_detail(sd){}

    friend void swap(X& lhs, X& rhs)
    {
        if(&lhs==&rhs)
            return;
        std::lock(lhs.m, rhs.m);
        std::lock_guard<std::mutex> lock_a(lhs.m, std::adopt_lock);
        std::lock_guard<std::mutex> lock_b(rhs.m, std::adopt_lock);
        swap(lhs.some_detail, rhs.some_detail);
    }
};
```

1

2
3

Alternative

```
void transfer(bank_account &from, bank_account &to, int amount)
{
    // lock both mutexes without deadlock
    std::lock(from.m, to.m);
    // make sure both already-locked mutexes are unlocked at the end of scope
    std::lock_guard<std::mutex> lock1(from.m, std::adopt_lock);
    std::lock_guard<std::mutex> lock2(to.m, std::adopt_lock);

    // equivalent approach:
    //     std::unique_lock<std::mutex> lock1(from.m, std::defer_lock);
    //     std::unique_lock<std::mutex> lock2(to.m, std::defer_lock);
    //     std::lock(lock1, lock2);

    from.balance -= amount;
    to.balance += amount;
}
```

Shared mutex for read only data

- Std::shared_mutex + std::shared_lock

```
class ThreadSafeCounter {
public:
    ThreadSafeCounter() = default;

    // Multiple threads/readers can read the counter's value at the same time.
    unsigned int get() const {
        std::shared_lock lock(mutex_);
        return value_;
    }

    // Only one thread/writer can increment/write the counter's value.
    void increment() {
        std::unique_lock lock(mutex_);
        value_++;
    }

    // Only one thread/writer can reset/write the counter's value.
    void reset() {
        std::unique_lock lock(mutex_);
        value_ = 0;
    }

private:
    mutable std::shared_mutex mutex_;
    unsigned int value_ = 0;
};
```

synchronization

- Conditional variable
- block a thread until another thread both modifies a shared variable (the *condition*), and notifies the condition_variable

[std::condition_variable](#)

Member functions

[condition_variable::condition_variable](#)
[condition_variable::~condition_variable](#)

Notification

[condition_variable::notify_one](#)
[condition_variable::notify_all](#)

Waiting

[condition_variable::wait](#)
[condition_variable::wait_for](#)
[condition_variable::wait_until](#)

Native handle

[condition_variable::native_handle](#)

```
#include <iostream>
#include <condition_variable>
#include <thread>
#include <chrono>

std::condition_variable_any cv;
std::mutex cv_m;
int i = 0;
bool done = false;

void waits()
{
    std::unique_lock<std::mutex> lk(cv_m);
    std::cout << "Waiting... \n";
    cv.wait(lk, [](){return i == 1;});
    std::cout << "...finished waiting. i == 1\n";
    done = true;
}

void signals()
{
    std::this_thread::sleep_for(std::chrono::seconds(1));
    std::cout << "Notifying...\n";
    cv.notify_one();

    std::unique_lock<std::mutex> lk(cv_m);
    i = 1;
    while (!done) {
        lk.unlock();
        std::this_thread::sleep_for(std::chrono::seconds(1));
        lk.lock();
        std::cerr << "Notifying again...\n";
        cv.notify_one();
    }
}

int main()
{
    std::thread t1(waits), t2(signals);
    t1.join(); t2.join();
}
```

Condition_variable

The thread that intends to modify the variable has to

1. acquire a `std::mutex` (typically via `std::lock_guard`)
2. perform the modification while the lock is held
3. execute `notify_one` or `notify_all` on the `std::condition_variable` (the lock does not need to be held for notification)

Even if the shared variable is atomic, it must be modified under the mutex in order to correctly publish the modification to the waiting thread.

Any thread that intends to wait on `std::condition_variable` has to

1. acquire a `std::unique_lock<std::mutex>`, on the same mutex as used to protect the shared variable
2. execute `wait`, `wait_for`, or `wait_until`. The wait operations atomically release the mutex and suspend the execution of the thread.
3. When the condition variable is notified, a timeout expires, or a `spurious wakeup` occurs, the thread is awakened, and the mutex is atomically reacquired. The thread should then check the condition and resume waiting if the wake up was spurious.

Spurious wakeup

- Is a hardware related issue

According to David R. Butenhof's Programming with POSIX Threads [ISBN 0-201-63392-2](#):

"This means that when you wait on a condition variable, the wait may (occasionally) return when no thread specifically broadcast or signaled that condition variable. Spurious wakeups may sound strange, but on some multiprocessor systems, making condition wakeup completely predictable might substantially slow all condition variable operations. The [race conditions](#) that cause spurious wakeups should be considered rare."

`wait` causes the current thread to block until the condition variable is notified or a spurious wakeup occurs, optionally looping until some predicate is satisfied.

- 1) Atomically unlocks lock, blocks the current executing thread, and adds it to the list of threads waiting on `*this`. The thread will be unblocked when `notify_all()` or `notify_one()` is executed. It may also be unblocked spuriously. When unblocked, regardless of the reason, lock is reacquired and wait exits. If this function exits via exception, lock is also reacquired. (until C++14)
- 2) Equivalent to

```
while (!pred()) {  
    wait(lock);  
}
```


`this_thread namespace`
`get_id(C++11)` `sleep_for(C++11)`
`yield(C++11)` `sleep_until(C++11)`

This_thread::sleep and yield

- They both causes the current thread to block
- sleep (until a condition)
- Yield: busy sleep
- Yield is better for fast turn-over model but waste CPU

```
void little_sleep(std::chrono::microseconds us)
{
    auto start = std::chrono::high_resolution_clock::now();
    auto end = start + us;
    do {
        std::this_thread::yield();
    } while (std::chrono::high_resolution_clock::now() < end);
}
```

Package task/promise/future

```
#include <iostream>
#include <future>
#include <thread>

int main()
{
    // future from a packaged_task
    std::packaged_task<int()> task([](){ return 7; }); // wrap the function
    std::future<int> f1 = task.get_future(); // get a future
    std::thread(std::move(task)).detach(); // launch on a thread

    // future from an async()
    std::future<int> f2 = std::async(std::launch::async, [](){ return 8; });

    // future from a promise
    std::promise<int> p;
    std::future<int> f3 = p.get_future();
    std::thread( [](std::promise<int>& p){ p.set_value(9); },
                std::ref(p) ).detach();

    std::cout << "Waiting..." << std::flush;
    f1.wait();
    f2.wait();
    f3.wait();
    std::cout << "Done!\nResults are: "
              << f1.get() << ' ' << f2.get() << ' ' << f3.get() << '\n';
}
```

async (lazy threading)

```
#include <iostream>
#include <vector>
#include <algorithm>
#include <numeric>
#include <future>

template <typename RAITer>
int parallel_sum(RAITer beg, RAITer end)
{
    typename RAITer::difference_type len = end-beg;
    if(len < 1000)
        return std::accumulate(beg, end, 0);

    RAITer mid = beg + len/2;
    auto handle = std::async(std::launch::async,
                           parallel_sum<RAIter>, mid, end);
    int sum = parallel_sum(beg, mid);
    return sum + handle.get();
}

int main()
{
    std::vector<int> v(10000, 1);
    std::cout << "The sum is " << parallel_sum(v.begin(), v.end()) << '\n';
}
```

Jthread (C++20)

- The class `jthread` represents [a single thread of execution](#). It has the same general behavior as [`std::thread`](#), except that `jthread` automatically rejoins on destruction, and can be cancelled/stopped in certain situations.

```
1 #include "jthread.hpp"
2 #include <chrono>
3 #include <iostream>
4
5 void sleep(const int seconds)
6 {
7     std::this_thread::sleep_for(std::chrono::seconds(seconds));
8 }
9
10 int main()
11 {
12     std::jthread jt{ [] (std::stop_token st) {
13         while (!st.requested_stop()) {
14             std::cout << "Doing work\n";
15             sleep(1);
16         }
17     }};
18     sleep(5);
19     jt.request_stop();
20     jt.join();
21 }
```

Coroutines: lightweight threads

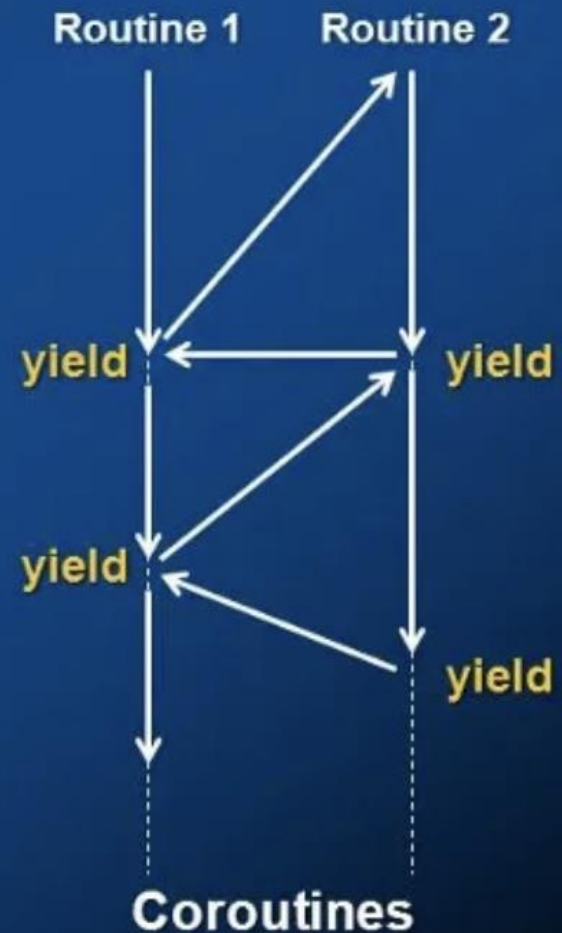
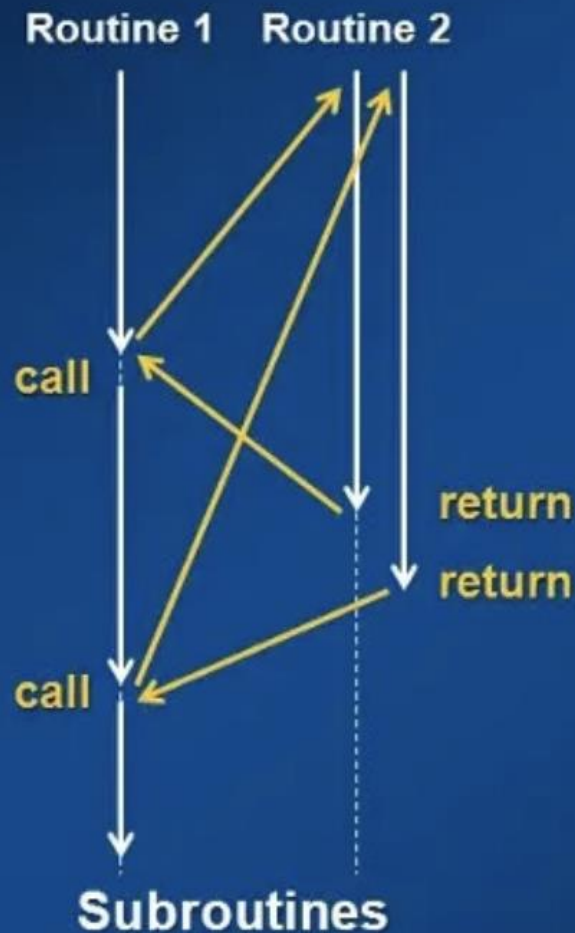
- **Coroutines** are stackless functions designed for enabling co-operative Multitasking, by allowing execution to be suspended and resumed.
 - **Cooperative Multitasking**—If multitasking participant process or thread voluntarily let go of control periodically or when idle or logically blocked.
- **Coroutines** suspend execution by returning to the caller and the data that is required to resume execution is stored separately from the stack. This allows for sequential code that executes asynchronously
 - well-suited for implementing familiar program components such as cooperative tasks, exceptions, event loops, state machines, and pipes.

Subroutines vs. Coroutines

"Subroutines are a special case of coroutines."

--Donald Knuth

Fundamental Algorithms. The Art of Computer Programming



Coroutines vs Threads

- **Coroutines provide concurrency but not parallelism**
- Switching between coroutines need not involve any system/blocking calls so no need for synchronization primitives such as mutexes, semaphores.
- Coroutines provide asynchronicity and resource locking isn't needed.

Coroutines in C++20

- A function is a coroutine if its definition does any of the following:
 - uses the `co_await` operator to suspend execution until resumed.
 - uses the keyword `co_yield` to suspend execution returning a value.
 - uses the keyword `co_return` to complete execution. Let's take a similar example to get a range.


```
#include <iostream>
```

```
#include <vector> // Coroutine gets called on need
```

```
generator <int> generateNumbers(int begin, int inc = 1) {
```

```
    for (int i = begin;; i += inc) {  
        co_yield i;  
    }
```

```
}
```

```
int main() {
```

```
    std::cout << std::endl;
```

```
    const auto numbers = generateNumbers(-10);
```

```
    for (int i = 1; i <= 20; ++i)  
        std::cout << numbers << " "; // Runs finite = 20 times***
```

```
    for (auto n:generateNumbers(0, 5)) // Runs infinite times**  
        std::cout << n << " "; // (3)
```

```
    std::cout << "\n\n";
```

```
}
```

CppCoro - A coroutine library for C++

The 'cppcoro' library provides a large set of general-purpose primitives for m proposal described in [N4680](#).

These include:

- Coroutine Types
 - `task<T>`
 - `shared_task<T>`
 - `generator<T>`
 - `recursive_generator<T>`
 - `async_generator<T>`

Atomic

```
//long multiThreadedSum = 0;  
std::atomic<long> multiThreadedSum( 0 );
```

```
long multiThreadedSum = 0;  
  
void SumNumbers( const std::vector<int>& toBeSummed, int idxStart, int idxEnd )  
{  
    for (int i = idxStart; i <= idxEnd; ++i)  
    {  
        multiThreadedSum += toBeSummed[i];  
    }  
}
```

```
int main()  
{  
    volatile int val = 0;  
    val += 1;  
    return val;  
}
```

```
main:  
    mov     DWORD PTR [rsp-4], 0  
    mov     eax, DWORD PTR [rsp-4]  
    add     eax, 1  
    mov     DWORD PTR [rsp-4], eax  
    mov     eax, DWORD PTR [rsp-4]  
    ret
```

Memory order

- Relaxed ordering
- Release-Acquire ordering
- Release-Consume ordering
- Sequentially-consistent ordering

```
1          int x = 0;      // global variable
2          int y = 0;      // global variable
3
4  Thread-1:                Thread-2:
5  x = 100;                 while (y != 200)
6  y = 200;                 ;
7                          std::cout << x;
```

```
1  Thread-1:
2  y = 200;
3  x = 100;
```

```
1          int x = 0;      // global variable
2
3  Thread-1:                Thread-2:
4  x = 100;      // A      std::cout << x;      // B
```

Relaxed ordering

- No control at all

```
void f()
{
    for (int n = 0; n < 1000; ++n) {
        cnt.fetch_add(1, std::memory_order_relaxed);
    }
}
```

Atomic operations library

The atomic library provides components for fine-grained programming. Each atomic operation is indivisible with respect to the object. Atomic objects are [free of data races](#).

Defined in header <atomic>

Atomic types

`atomic` (C++11)

`atomic_ref` (C++20)

Operations on atomic types

`atomic_is_lock_free` (C++11)

`atomic_store` (C++11)

`atomic_store_explicit` (C++11)

`atomic_load` (C++11)

`atomic_load_explicit` (C++11)

`atomic_exchange` (C++11)

`atomic_exchange_explicit` (C++11)

`atomic_compare_exchange_weak` (C++11)

`atomic_compare_exchange_weak_explicit` (C++11)

`atomic_compare_exchange_strong` (C++11)

`atomic_compare_exchange_strong_explicit` (C++11)

`atomic_fetch_add` (C++11)

`atomic_fetch_add_explicit` (C++11)

`atomic_fetch_sub` (C++11)

`atomic_fetch_sub_explicit` (C++11)

`atomic_fetch_and` (C++11)

`atomic_fetch_and_explicit` (C++11)

`atomic_fetch_or` (C++11)

`atomic_fetch_or_explicit` (C++11)

Release-Acquire ordering

memory_order_acquire	A load operation with this memory order performs the <i>acquire operation</i> on the affected memory location: no reads or writes in the current thread can be reordered before this load. All writes in other threads that release the same atomic variable are visible in the current thread (see Release-Acquire ordering below)
memory_order_release	A store operation with this memory order performs the <i>release operation</i> : no reads or writes in the current thread can be reordered after this store. All writes in the current thread are visible in other threads that acquire the same atomic variable (see Release-Acquire ordering below) and writes that carry a dependency into the atomic variable become visible in other threads that consume the same atomic (see Release-Consume ordering below).

```
std::atomic<bool> ready{ false };
int data = 0;

void producer()
{
    data = 100; // A
    ready.store(true, std::memory_order_release); // B
}

void consumer()
{
    while (!ready.load(std::memory_order_acquire)) // C
        ;
    assert(data == 100); // never failed // D
}
```