

# Multiprocessor Real-Time Scheduling

Real-Time and Embedded Operating  
Systems

Prof. Li-Pin Chang  
ESSLab@NYCU

# Outline

- Multiprocessor Real-Time Scheduling
- Global Scheduling
- Partitioned Scheduling
- Semi-partitioned Scheduling

# Multiprocessor Models

- Identical (Homogeneous) processors:
  - All processors are made of the same hardware
  - All processors have the same clock rate
  - This unit
- Uniform processors
  - All processors are made of the same hardware
  - Processors have different clock rates
  - A job runs faster on fast-clocked processors
  - DVFS

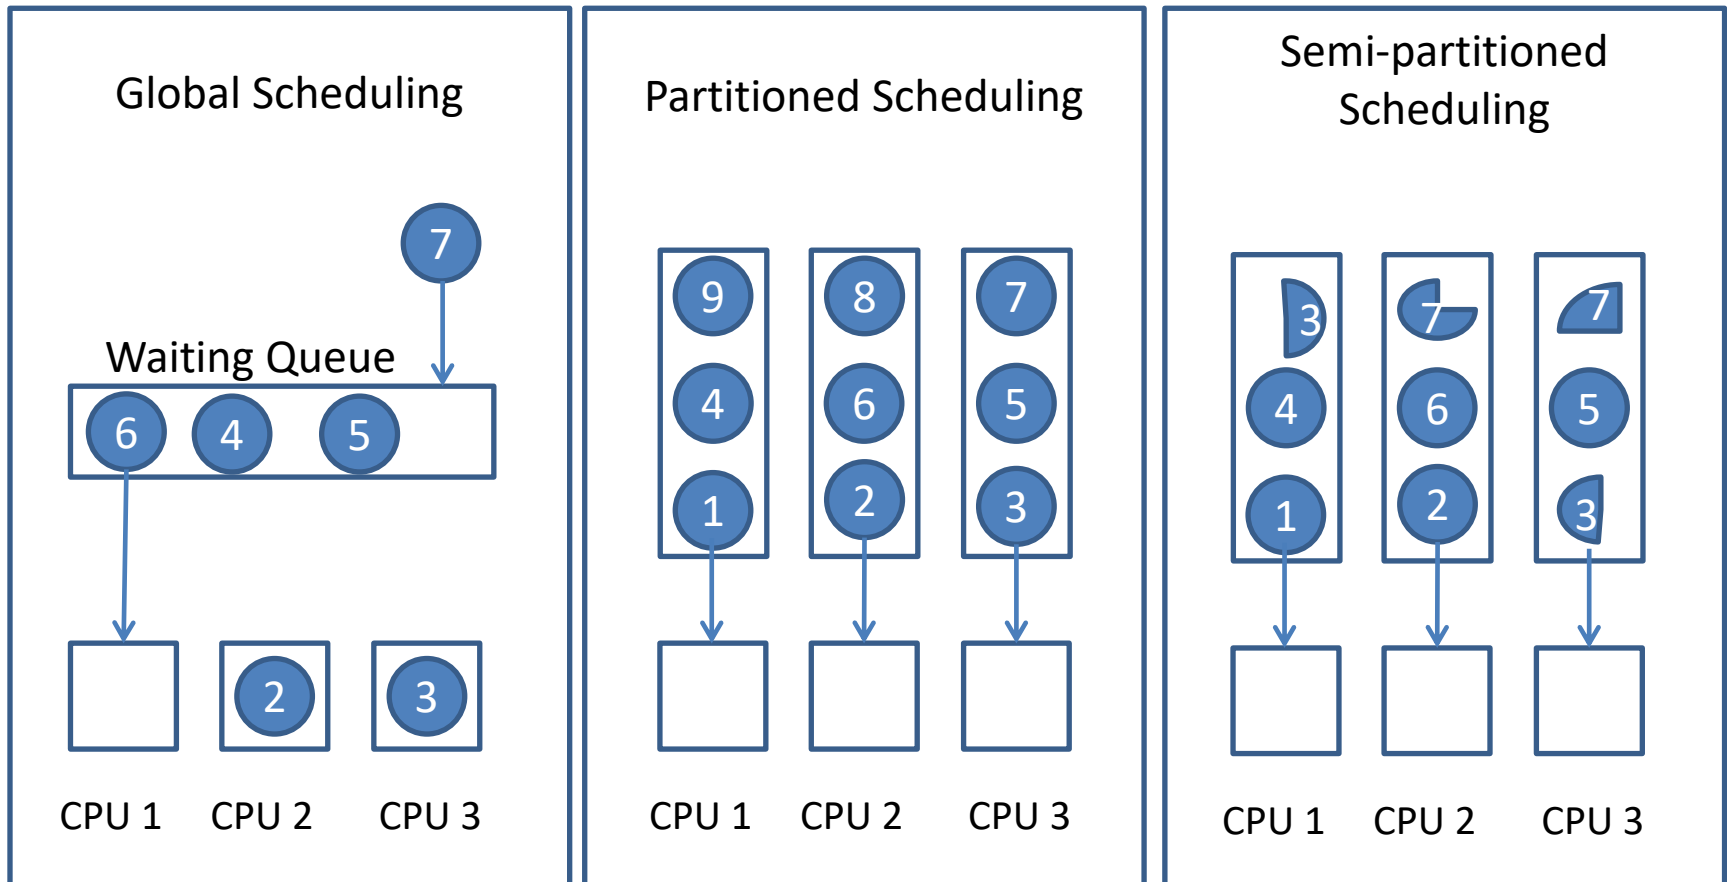
# Multiprocessor Models

- Unrelated (Heterogeneous) processors
  - A job has different execution times on different processors
  - A processor may execute a job faster than other processors but execute another job slower than other processors
  - For example, multiprocessors with different instruction set architectures (ISAs)

# Scheduling Models

- Global Scheduling:
  - A job can be dispatched to any processor
  - Job migrate among processors whenever necessary
  - A global ready queue
- Partitioned Scheduling:
  - Tasks are statically partitioned among processors
  - No task migration is allowed
  - Per-processor ready queues
- Semi-partitioned Scheduling:
  - Based on partitioned scheduling
  - Involves limited on-line job migration

# Scheduling Models



# Global Scheduling

- Here, a ready task/job means a task that can be executed
- It can be
  - A task waits in the ready queue
  - A task is being executed on a processor

# Global Scheduling

- All ready jobs are kept in a global queue, and a job can be migrated to any processor
- **Global-EDF:** When a job finishes or a new job arrives at the global queue, the  $M$  processor executes  $M$  ready jobs having the  $M$  shortest deadlines
- **Global-RM:** When a job finishes or a new job arrives at the global queue, the  $M$  processor executes  $M$  ready jobs having the  $M$  shortest periods




# Global Scheduling

- The  $M$  processors always execute  $M$  ready jobs with the  $M$  earliest deadlines
- When a new job arrives
  1. If there is an idle processor, use it
  2. Otherwise, if it can preempt a running job, it preempts the running job having the farthest deadline
- To avoid shuffling tasks on processors

# Global EDF

- $\{t1=(2,3), t2=(2,4), t3=(8,12)\}$

task set: schedulable



# Global Scheduling

- Advantages:
  - Good processor utilization
  - Unused processor time can easily be reclaimed during run-time for soft RT tasks
- Disadvantages:
  - Less intuitive! Many single-processor scheduling results cannot be extended to multiprocessor global scheduling
  - Adding processors, reducing task computation times, or “enhance” other system parameters can unexpectedly degrade task response!


# Scheduling Anomaly

- Increasing the period of a task may negatively impact on the response time of another task

# Anomaly 1: Relaxing Task Period


task set: schedulable

$\{t1=(2,3), t2=(2,4), t3=(8,12)\}$



task set: unschedulable

$\{t1=(2,4), t2=(2,4), t3=(8,12)\}$



$\tau_3$  needs to execute  
two more time units.

# Anomaly 2: Dhall's Effect

- Dhall's Effect of global scheduling

| Task | P  | C | U    |
|------|----|---|------|
| T1   | 10 | 5 | 0.5  |
| T2   | 10 | 5 | 0.5  |
| T3   | 12 | 8 | 0.67 |

- T3 is not schedulable by global EDF/RM
  - but is schedulable if T1 and T2 share the same processor

# Schedulability Test

- A set of periodic tasks  $t_1, t_2, \dots, t_N$  with implicit deadlines is schedulable on  $M$  processors using preemptive Global EDF scheduling if

$$\sum_{i=1}^N \frac{C_i}{T_i} \leq M(1 - \frac{C_k}{T_k}) + \frac{C_k}{T_k},$$

where  $t_k$  is the task of the largest utilization  $C_k/T_k$

# Weakness of Global Scheduling

- Scheduling Anomaly
- Migration overhead
  - Cache re-population (cold start)
  - Pipeline stall



# Partitioned Scheduling

- Two steps:
  - Partitioning tasks among processors
  - Scheduling tasks on each processor
- Example: Partitioned scheduling with EDF
  - Assign tasks to the processors such that no processor's capacity exceeds 100%
  - Schedule tasks on each processor using EDF

# Partitioned Scheduling

- Advantages:
  - Most techniques for single-processor scheduling are applicable here
- Partitioning of tasks can be automated
  - Solving a bin-packing problem
- Disadvantages:
  - Cannot reclaim unused processor time
  - May have very low utilization, bounded by 50%
    - Worst case of bin-packing heuristic

# Task Partitioning Problem

Given a set of tasks with arbitrary deadlines, the objective is to find a feasible task assignment onto  $M$  processors such that all the tasks meet their timing constraints

# Bin Packing Problem

## Optimization version

Given a bin size  $V$  and a list  $a_1, \dots, a_n$  of sizes of the items to pack, find an integer  $B$  and a  $B$ -partition  $S_1 \cup \dots \cup S_B$  of  $\{1, \dots, n\}$  such that  $\sum_{i \in S_k} a_i \leq V$ , for all  $k = 1, \dots, B$ .  
A solution is optimal if it has minimal  $B$ .

## Decision version

The same as above, but asking whether all the objects can be packed into  $B$  bins of the same capacity.

The decision version of Bin Packing is known to be NP-complete, which can be reduced (transformed) to an instance of partitioned scheduling.

# Bin-Packing versus Partitioned Scheduling


- Bin Packing: packing objects of varying sizes in boxes ("bins") with the objective of minimizing number of used boxes.
  - Solutions (Heuristics): First Fit, etc.
- Application to multiprocessor systems:
  - Bins are represented by processors and objects by tasks
  - The decision whether a processor is "full" or not is derived from a utilization-based schedulability test

# Partitioning Algorithms

- First-Fit: choose the fitting processor of the smallest index
- Best-Fit: choose the fitting processor of the maximal utilization
- Worst-Fit: choose the fitting processor of the minimal utilization

# Partitioned Example

- 0.2 -> 0.6 -> 0.4 -> 0.7 -> 0.1 -> 0.3



First Fit  
(not Next Fit)



Best Fit



Worst Fit

# Schedulability Test

Lopez [3] proves that the **worst-case achievable utilization** for EDF scheduling and FF allocation (EDF-FF) takes the value

If all the tasks have an utilization factor  $C/T$  under a value  $\alpha$ , where  $m$  is the number of processors

$$U_{wc}^{EDF-FF}(m, \beta) = \frac{\beta m + 1}{\beta + 1} \quad \text{where} \quad \beta = \lfloor 1 / \alpha \rfloor$$



# Weakness of Partitioned Scheduling


- Binding a task to a processor makes the problem NP-hard and cause pessimistic results
- Example: Suppose that there are  $M$  processors and  $M + 1$  tasks with the same period  $T$  and the (worst-case) execution times of all these  $M + 1$  tasks are  $T/2 + e$  with  $e > 0$ 
  - With partitioned scheduling, it is not schedulable
  - Is it possible to divide a task between two processors?

# Semi-partitioned Scheduling

- Based on First Fit
- Adding tasks to a processor until the processor is fully loaded
- Partitioning the next task into  $p_1$  and  $p_2$  and completely fill the current processor with  $p_1$
- Adding  $p_2$  to the next processor

# Semi-partitioned EDF

- Assignment phase
  - Applying first-fit algorithm, by taking **SEP** as the upper bound of utilization on a processor.
  - If a task does not fit, split this task into two subtasks, one is assigned on the current processor and the other is assigned to the next processor



# Semi-partitioned EDF


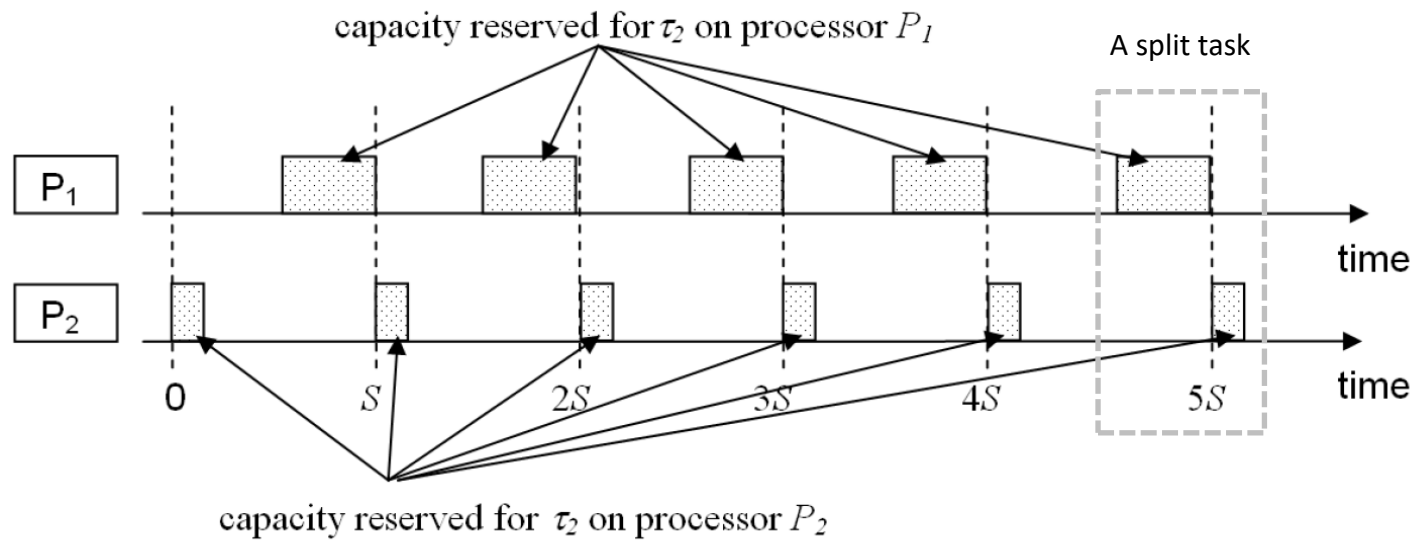
- We can assign all the tasks  $t_i$  with  $U_i > SEP$  on a dedicated processor. So, we only consider tasks with  $U_i$  no larger  $SEP$ .

```
1:  $m \leftarrow 1, U_m \leftarrow 0;$ 
2: for  $i = 1$  to  $N$ , where  $N = |\mathbf{T}|$  do
3:   if  $\frac{C_i}{T_i} + U_m \leq SEP$  then
4:     assign task  $\tau_i$  on processor  $m$ ;
5:      $U_m \leftarrow U_m + \frac{C_i}{T_i};$ 
6:   else
7:     assign task  $\tau_i$  on processor  $m$  with  $lo\_split(\tau_i)$  set to  $SEP - U_m$  and on
       processor  $m + 1$  with  $high\_split(\tau_i)$  set to  $\frac{C_i}{T_i} - (SEP - U_m);$ 
8:      $m \leftarrow m + 1$  and  $U_m \leftarrow \frac{C_i}{T_i} - (SEP - U_m);$ 
```

When executing, the reservation to serve  $t_i$  is to set  $x_i$  to  $S \times (f + lo\_split(t_i))$  and  $y_i$  to  $S \times (f + high\_split(t_i))$ .  $SEP$  is set as a constant.

# Semi-partitioned EDF

- Execution phase
  - $T_{\min}$  is the minimum period among all the tasks
  - By a user-designed parameter  $k$ , we divide time into slots with length  $S = T_{\min}/k$
  - Execution of a split task is only possible in the reserved time window in the time slot
  - The rest of the time: scheduling tasks on each individual processor using EDF




 Execution windows of non-split job




Execution window of split job

# Semi-partitioned EDF

- For each time slot, we will reserve two parts.



If a task  $t_i$  is split, the task can be served only within these two pre-defined time slots with length  $x_i$  and  $y_i$ .



A processor can host two split tasks,  $t_i$  and  $t_j$ .  $t_i$  is served at the beginning of the time slot, and  $t_j$  is served at the end.

The schedule is EDF, but if a split task instance is in the ready queue, it is executed in the reserved time region.

# Two Split Tasks on a Processor

- For split tasks to be schedulable, the following sufficient conditions have to be satisfied
  - $\text{lo\_split}(t_i) + f + \text{high\_split}(t_i) + f \leq 1$  for any split task  $t_i$ .
  - $\text{lo\_split}(t_j) + f + \text{high\_split}(t_i) + f \leq 1$  when  $t_i$  and  $t_j$  are assigned on the same processor.
- Therefore, the “magic value” SEP

$$SEP \leq 1 - 2f \leq 1 - 2(\sqrt[2]{\kappa(\kappa + 1)} - \kappa).$$

- However, we still have to guarantee the schedulability of the non-split tasks. It can be shown that the sufficient condition is

$$SEP \leq 1 - 4f \leq 1 - 4(\sqrt[2]{\kappa(\kappa + 1)} - \kappa).$$

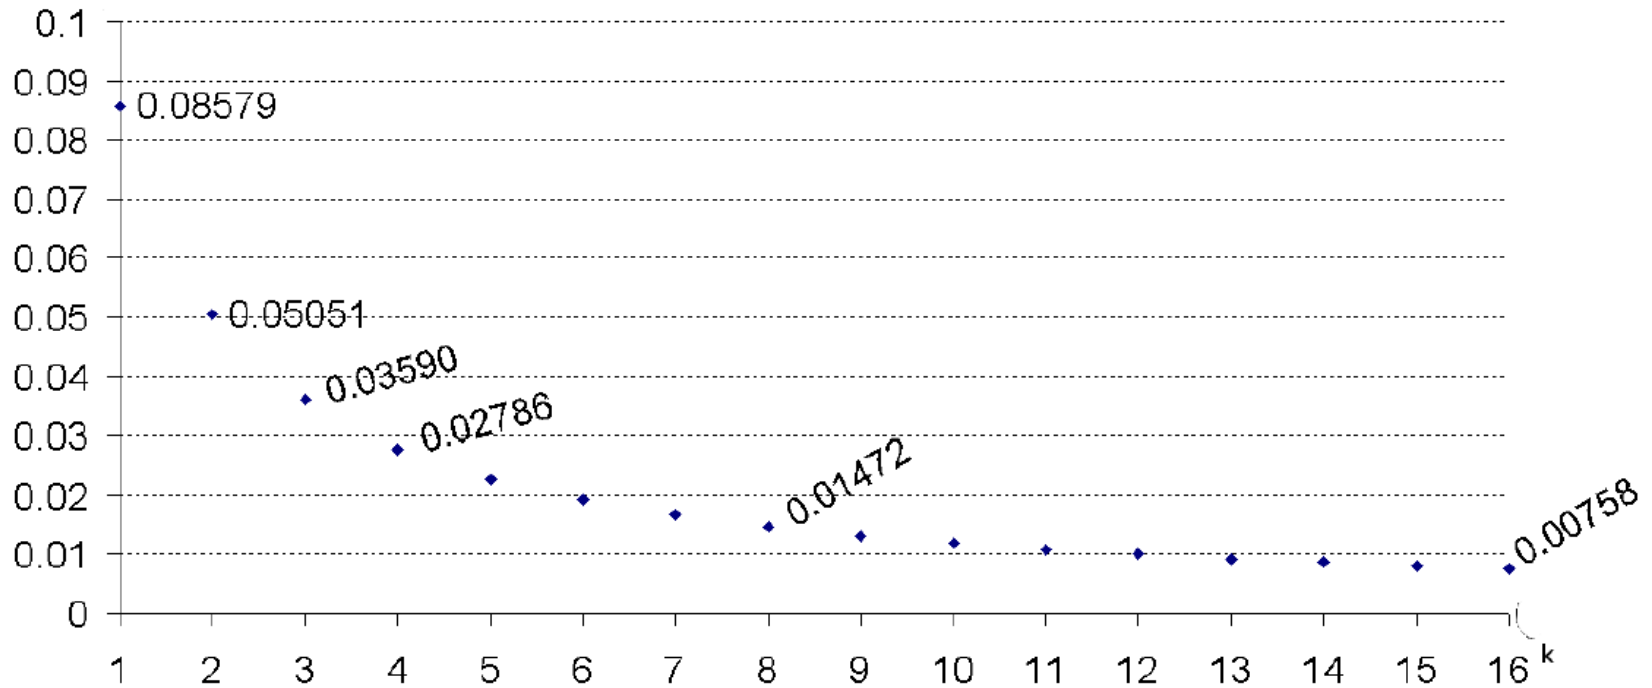


# Schedulability Test

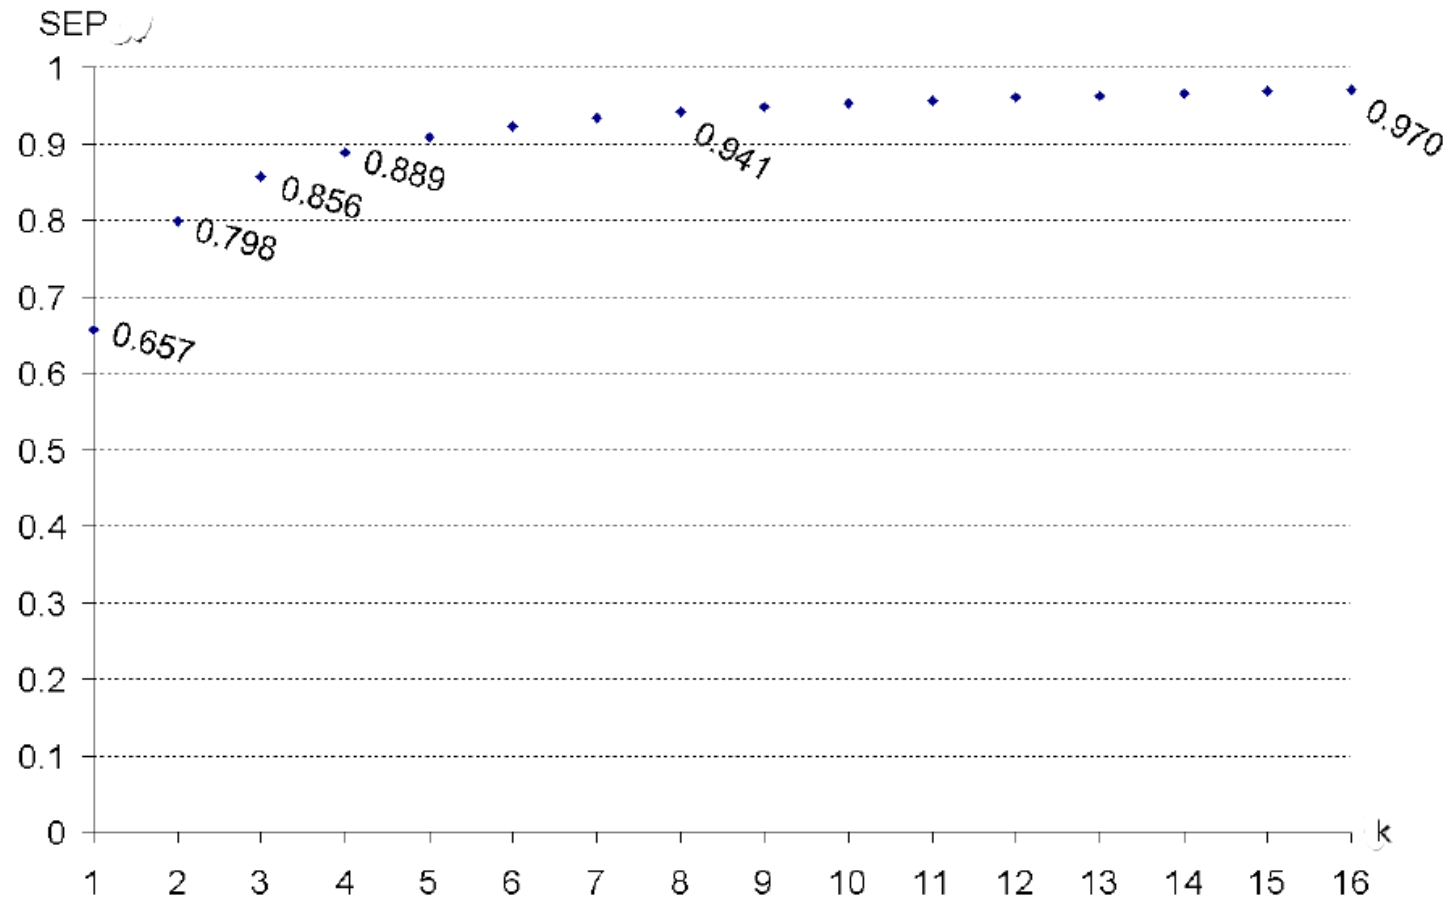
By taking  $SEP$  as  $1 - 4(\sqrt[2]{\kappa(\kappa + 1)} - \kappa)$  and  $f = \sqrt[2]{\kappa(\kappa + 1)} - \kappa$ , the above algorithm guarantees to derive feasible schedule if  $\sum_{\tau_i \in \mathbf{T}} \frac{C_i}{T_i} \leq M' \cdot SEP$  and  $\frac{C_i}{T_i} \leq SEP$  for all tasks  $\tau_i$ .

$M'$  = the # of processors serving tasks whose individual utilization  $\leq SEP$

# Magic Values: f



# Magic Values: SEP



# Reference

- Multiprocessor Real-Time Scheduling
  - Dr. Jian-Jia Chen: Multiprocessor Scheduling. Karlsruhe Institute of Technology (KIT): 2011-2012
- Global Scheduling
  - Sanjoy K. Baruah: Techniques for Multiprocessor Global Schedulability Analysis. RTSS 2007: 119-128
- Partitioned Scheduling
  - Sanjoy K. Baruah, Nathan Fisher: The Partitioned Multiprocessor Scheduling of Sporadic Task Systems. RTSS 2005: 321-329
- Semi-partitioned Scheduling
  - Björn Andersson, Konstantinos Bletsas: Sporadic Multiprocessor Scheduling with Few Preemptions. ECRTS 2008: 243-252

# Credit

- This slice set is based on materials provided by Prof. Ya-Shu Chen (NTUST)