

Resource Synchronization Protocols

Real-Time and Embedded Operating
Systems

Prof. Li-Pin Chang
ESSLab@NYCU

Motivation

- Tasks share and compete for resources
- Resource access from tasks must be synchronized
 - To avoid race conditions
 - To avoid deadlocks
 - To manage priority inversion
- Acquisition (lock) -> use -> release
 - Acquisition is *never* rejected but may be postponed
 - The waiting time must be predictable

Okay, but first...

- What kind of system does not need synchronization on resource access?
 - Non-preemptible scheduling
 - A job releases all resources upon its completion
 - Cyclic executive or frame-based scheduling
 - Resource access is statically programmed

Resources

- Let us consider **passive** resources only
- When a task is using a passive resource, it is busy on the CPU. It does not suspend on the CPU.
 - E.g., data structures protected by sems or mutexes
- Active resources are not considered here
 - A task suspends on the CPU when using an active resource, e.g., I/O devices

Resources

- Consider n types of serially reusable resources R_1, R_2, \dots, R_n
- R_i has v_i indistinguishable units of resource (instances)
 - For example
 - A global buffer has 20 empty slots
 - A DMA controller has 4 channels
 - Acquiring which one of the instances does not matter

Resources

- A job **locks** a resource before using it
 - The lock attempt may not be granted immediately
 - The task **waits** until the lock attempt is granted
- Locking is not granted immediately (and is postponed) bcz
 - the resource is not currently available, or granting the lock may cause some undesirable effects
- When a job no longer needs a resource, it **unlocks** it
 - Unlocking is effective immediately
- Both locking and unlocking may cause context switch

Resources

- Let the duration between the locking and unlocking of a resource be a **critical section**
- Let all critical sections be **properly nested**




FIGURE 8-1 Examples of critical sections

Resources






FIGURE 8-2 Example of job interaction due to resource contention.

- Resource contention increases task response time
 - If J_1 and J_2 do not require resources, they can finish by time 11 and 14, respectively (vs 12 and 17)

Resources

- A job is directly blocked if the requested resource is held by another job
- Use a wait-for graph to describe the run-time resource usage
- Use a bipartite graph to describe resource requirements



Undesirable Effects - I

- Priority inversion (blocking)
 - A high-priority job can not run because of resource contention from low-priority jobs
 - E.g., the high-priority job tries to lock a resource that is currently locked by a low-priority job
 - In this case, the high-priority job is **blocked**
 - Note: a low-priority job is never “blocked” by high-priority jobs (why?)
 - May damage the schedulability of HPTs

What happened on Mars?



← Pathfinder,
protected by airbags

→ The Pathfinder



**Pictures courtesy of NASA

Pathfinder Mars Landing



Undesirable Effects – I (cont'd)

- Even worse a, job may be *indirectly* blocked by another job, even if the two jobs do not share any resources!!
 - uncontrollable priority inversion




FIGURE 8-4 Uncontrolled priority inversion.

On Pathfinder, this medium-priority job cause a time-out of the high-priority J_1 !

Undesirable Effects - II

- Timing anomalies
 - Ideally, if a job completes earlier than expected, then the response time of all the other jobs become earlier (at least, not later)
 - But exceptions exist..




FIGURE 8-2 Example of job interaction due to resource contention.





FIGURE 8-3 Example illustrating timing anomaly.

Now the duration of this CS becomes shorter...

Undesirable Effects - III

- Deadlocks
 - A **fatal error** in real-time systems
 - If a deadlock occurs, then some jobs are involved in circular waiting
 - Deadlock prevention: enforcing a partial order on resource access



Undesirable Effects

- Priority inversions and timing anomalies are inevitable with mutual-exclusive access
 - They should be **managed** (e.g., **bounded**)
- Deadlocks are a fatal error
 - They must **be avoided**

Resource Synchronization Protocols

- Protocol: a set of rules
 - to grant and to postpone lock requests
 - to schedule jobs (priority manipulation)

Non-Preemptible Critical Section

- **NPCS**: A job becomes non-preemptible when it is using a resource
 - Any HPT cannot preempt it
 - Equivalent to locking the scheduler

Non-Preemptible Critical Section




FIGURE 8-2 Example of job interaction due to resource contention.

← Ungoverned
(mutual exclusion only)

NPCS→



FIGURE 8-2 Example of job interaction due to resource contention.

Non-Preemptible Critical Section

- **Deadlocks** never occur
 - A job holding a resource can never be blocked (no hold and wait)
 - (why?)
- **Uncontrollable priority inversions** never occur
 - (why?)

Non-Preemptible Critical Section

- Let tasks in $\{T_1, T_2, \dots, T_n\}$ be sorted in the rate-monotonic order and scheduled by **RM**
- The longest blocking time imposed on task T_i because of resource contention is

$$\max_{i+1 \leq k \leq n} (s_k)$$

- s_k stands for the longest critical section of task T_k

Why at most once?

Non-Preemptible Critical Section

- If tasks are scheduled by **EDF**, the longest blocking time of T_i is

$$\max_{i+1 \leq k \leq n} (s_k)$$

- Still the same (huh?)

Schedulability tests for RM/EDF with NPCS will be discussed later


Non-Preemptible Critical Section

- Pros:
 - Easy to implement
 - No need to know resource usage *a priori*
- Cons:
 - Poor response (high priority tasks suffer)
 - An LPT blocks a HPT even if they do not share any resources

Ceiling-Priority Protocol

- Used in Ada real-time programming language
 - Commonly used in aerospace or aviation systems
- Resource usages are known *a priori*
- Consider a task set $T = \{T_a, T_b, T_c, T_d, T_e, T_f, \dots\}$, sorted by priority (high- \rightarrow low)
 - Let T_b, T_d, T_e , share a resource R
 - When a task T locks R , its priority is boosted to T_b 's priority until it unlocks R

Ceiling-Priority Protocol




Ceiling-Priority Protocol

- As known as “highest-locker protocol”
- Better response than NPCS
- Free from
 - Uncontrollable priority inversion
 - Deadlocks
 - These properties directly follow from NPCS

Ceiling-Priority Protocol

- Blocking time
- Consider T_a , T_b , T_c , and T_d
 - T_b and T_d share a resource R
 - T_a and T_c do not use any resources
 - the longest critical sections $c_b=1$ and $c_d=2$
 - Blocking time $b_a=0, b_b=2, b_c=2, b_d=0$
 - NPCS?

What Happened to Ariane 5?

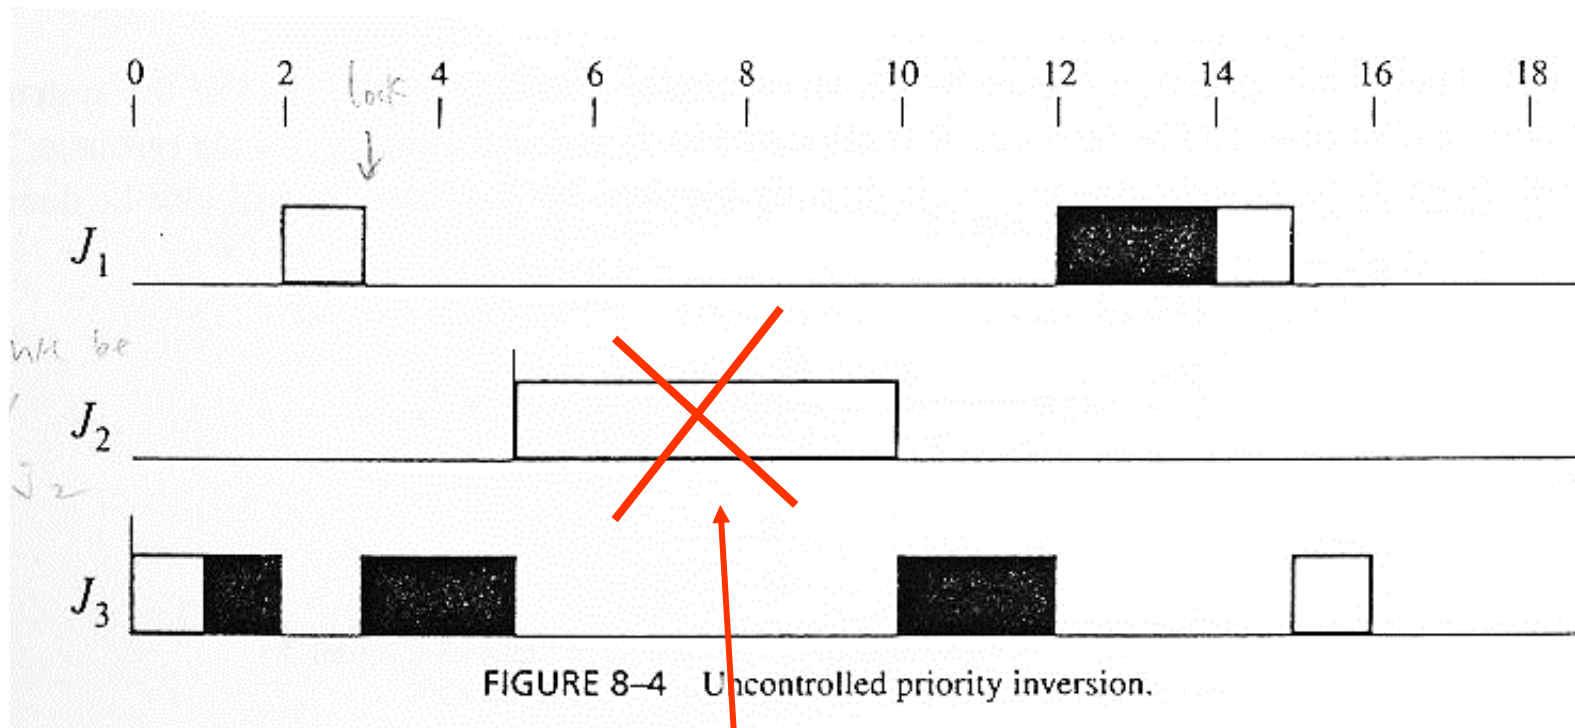


It's nothing to do with Ada. It's a software bug.

Priority-Inheritance Protocol

- NPCS is **too** restrictive because an LPT affects all tasks on locking
 - CPP suffer from the same problem but in a lower degree (highest locker's priority)
- PIP relieves such restriction
 - uncontrollable blocking must not occur
 - Priority is boosted on **blocking**, not locking

Priority-Inheritance Protocol



This job should not run because a higher-priority job is already waiting (being blocked)!!





FIGURE 8-4 Uncontrolled priority inversion.



(b)

When blocking J_1 , J_3 runs on behalf of J_1

Medium-priority job is not executed because there is a blocked job with a higher priority

Priority-Inheritance Protocol


Rules of the Basic Priority-Inheritance Protocol

1. *Scheduling Rule*: Ready jobs are scheduled on the processor preemptively in a priority-driven manner according to their current priorities. At its release time t , the current priority $\pi(t)$ of every job J is equal to its assigned priority. The job remains at this priority except under the condition stated in rule 3.
2. *Allocation Rule*: When a job J requests a resource R at time t ,
 - (a) if R is free, R is allocated to J until J releases the resource, and
 - (b) if R is not free, the request is denied and J is blocked.
3. *Priority-Inheritance Rule*: When the requesting job J becomes blocked, the job J_l which blocks J inherits the current priority $\pi(t)$ of J . The job J_l executes at its inherited priority $\pi(t)$ until it releases R ; at that time, the priority of J_l returns to its priority $\pi_l(t')$ at the time t' when it acquires the resource R .


Note: a job may inherits priority from one or more jobs (nested inheritance)

| Job | r_i | e_i | π_i | Critical Sections |
|-------|-------|-------|---------|--------------------------|
| J_1 | 7 | 3 | 1 | [Shaded; 1] |
| J_2 | 5 | 3 | 2 | [Black; 1] |
| J_3 | 4 | 2 | 3 | |
| J_4 | 2 | 6 | 4 | [Shaded; 4 [Black; 1.5]] |
| J_5 | 0 | 6 | 5 | [Black; 4] |

(a)




Transitive priority inheritance & Transitive blocking



Priority-Inheritance Protocol


- Chain blocking




- There can be multiple undergoing critical sections under PIP
- There is only one undergoing critical section under NPCS and CPP

Priority-Inheritance Protocol

- Both transitive and chain blocking could occur in PIP




Priority-Inheritance Protocol



J_1 can be blocked up to $\min(v, k)$ times, each of the duration of the outmost CS. v and k stand for different resources J_1 requires and the number of low-priority tasks, respectively

Priority-Inheritance Protocol

- Deadlocks



PIP allows hold and wait

Priority-Inheritance Protocol

- Pros:
 - Simple
 - Better response than NPCS and CPP
 - Free from the uncontrollable priority
- Cons:
 - Suffering from deadlocks
 - Chain/transitive blocking


Priority-Ceiling Protocol

- PCP (fixes) improves upon PIP in terms of
 - Prohibiting chain/transitive blocking
 - To reduce the worst-case blocking time
 - Prohibiting circular waiting
 - To avoid deadlock

Priority-Ceiling Protocol

- Assumptions
 - All tasks have unique and fixed priorities
 - Resource usages of tasks are known *a priori*
- Terms
 - $\Pi(R)$: priority ceiling of resource R
 - The highest priority of all tasks that require R
 - $\Pi^{\wedge}(t)$: current system ceiling at time t (*caution!*)
 - The highest among the ceilings of all the resources that are *currently in use*

Demonstration of System Ceiling



Priority-Ceiling Protocol

Rules of Basic Priority-Ceiling Protocol

1. Scheduling Rule:


- (a) At its release time r , the current priority $\pi(r)$ of every job J is equal to its assigned priority. The job remains at this priority except under the condition stated in rule 3.
- (b) Every ready job J is scheduled preemptively and in a priority-driven manner at its current priority $\pi(t)$.

2. Allocation Rule: Whenever a job J requests a resource R at time t , one of the following two conditions occurs:

- (a) R is held by another job. J 's request fails and J becomes blocked.
- (b) R is free.
 - (i) If J 's priority $\pi(t)$ is higher than the current priority ceiling $\hat{\Pi}(t)$, R is allocated to J .
 - (ii) If J 's priority $\pi(t)$ is not higher than the ceiling $\hat{\Pi}(t)$ of the system, R is allocated to J only if J is the job holding the resource(s) whose priority ceiling is equal to $\hat{\Pi}(t)$; otherwise, J 's request is denied, and J becomes blocked.

3. Priority-Inheritance Rule: When J becomes blocked, the job J_i which blocks J inherits the current priority $\pi(t)$ of J . J_i executes at its inherited priority until the time when it releases every resource whose priority ceiling is equal to or higher than $\pi(t)$; at that time, the priority of J_i returns to its priority $\pi_i(t')$ at the time t' when it was granted the resource(s).

Priority-Ceiling Protocol




Priority-Ceiling Protocol

- PCP may postpone a resource request even if the resource is currently available
 - To prevent undesirable effects, specifically, deadlocks and chain/transitive blocking
- There are three types of blockings in PCP
 - Direct blocking
 - Priority-inheritance blocking
 - Avoidance/ceiling blocking (new)

Priority-Ceiling Protocol



(a) Direct blocking



(b) Priority-inheritance blocking



J's current priority is lower than the ceiling of X

(c) Avoidance blocking

Priority-Ceiling Protocol

- PCP avoids uncontrolled blocking
 - Through priority inheritance
- PCP also avoids deadlocks





FIGURE 8-12 Example illustrating how priority-ceiling protocol prevents deadlock.

Priority-Ceiling Protocol


- PCP avoids deadlocks
(informal) Proof:



Jx and Jy can not
simultaneously lock
resources!!

Priority-Ceiling Protocol

- Theorem: PCP avoids deadlocks
(informal) Proof:



If any one job successfully locks some resources, then at least one of the other jobs can not lock resources

Priority-Ceiling Protocol

- Jobs seem having longer response under PCP than under PIP
 - PCP has priority-ceiling blocking but PIP doesn't
- In fact, the worse-case blocking time of PCP is not longer than PIP
 - No transitive blocking and chain blocking in PCP, tasks are blocked by up to *one* critical section

Priority-Ceiling Protocol


- Theorem: any job governed by PCP can be blocked for at most one critical section

Proof:

- A job can be blocked by only once (no chain)
- No transitive blocking


Priority-Ceiling Protocol

- A job can be blocked **only once**
 - When HPT acquires R, no other tasks $< \text{HPT}$ could hold a resource that HPT needs. Otherwise, HPT cannot acquire R (due to system ceiling)



Priority-Ceiling Protocol

- No transitive blocking
 - If JL locks R1 then J cannot lock R2



$J_H \rightarrow J \rightarrow J_L$

Priority-Ceiling Protocol

- Priority inheritance blocking and priority ceiling blocking





FIGURE 8-14 Example on duration of blocking.

- J_4 can block J_3 by inheriting a priority from J_1
 - J_4 can block J_3 by raising the system ceiling to $\pi(\text{Black})$
- Different scenarios, but the durations are the same

Priority-Ceiling Protocol



6,6,5,4,4

| | Directly blocked by | | | | | Priority-inher blocked by | | | | | Priority-ceiling blocked by | | | | |
|-------|---------------------|-------|-------|-------|-------|---------------------------|-------|-------|-------|-------|-----------------------------|-------|-------|-------|-------|
| | J_2 | J_3 | J_4 | J_5 | J_6 | J_2 | J_3 | J_4 | J_5 | J_6 | J_2 | J_3 | J_4 | J_5 | J_6 |
| J_1 | | 6 | | | 2 | | | | | | | | | | |
| J_2 | * | | 5 | | | * | 6 | | | 2 | * | 6 | | | 2 |
| J_3 | | * | | | 4 | | * | 5 | | 2 | | * | 5 | | 2 |
| J_4 | | | * | | | | | * | | 4 | | | * | | 4 |
| J_5 | | | | * | | | | | * | 4 | | | | * | |

max(2,4)

Priority-Ceiling Protocol

- For any job J **that requires some resource(s)**, its priority-inheritance blocking time and its ceiling blocking time are **the same**
 - PI & PC blockings happen through the same path in the resource allocation graph
 - But remember: ceiling is to manage “locking”, priority is to manage “scheduling”
 - So PI & PC blocking times will be **different** for tasks **not using any resource(s)**

Priority-Ceiling Protocol



- A collection of tasks $\{T_1, T_2, \dots, T_n\}$ governed by PCP are schedulable if

$$\forall_{i=1 \dots n} \left\{ \frac{b_i}{p_i} + \sum_{j=1}^i \frac{c_j}{p_j} \leq U(i) \right\}$$

- Where b_i is the longest blocking time imposed on T_i

Priority-Ceiling Protocol

$$\forall_{i=1 \dots n} \left\{ \frac{b_i}{p_i} + \sum_{j=1}^i \frac{c_j}{p_j} \leq U(i) \right\}$$



The critical instant of T_i (left) with b_i included in c_i , and (right) with blocking time

- With task blocking time determined, we can also use response time analysis
- What are the tests for NPCS and CPP?
 - The same tests (U bound and rsp time analysis) can be used, as “blocking time” remains the same regardless of the synchronization protocols

Applicability

- So far we are focused on RM, and the following protocols are applicable
 - NPCS, CPP, PIP, and PCP
- For EDF, we previously mentioned that NPCS and CPP can be used as well
 - EDF does not work with PIP and PCP, however
 - A similar but different approach, SRP, is to be discussed

Stack-Resource Policy

- PCP is not directly applicable to EDF, because with EDF there is no task priority, so
 - it is not possible to define ceilings
 - priority inheritance cannot operate
- In EDF, a task preempts longer-period tasks
- So “preemption levels” can be defined to be inversely proportional to task periods
 - It replaces the notion of “task priority”

Stack-Resource Policy

- Resource/system ceiling
 - Use preemption levels
- Inheritance
 - Use deadlines
- Scheduling (!)
 - A task cannot execute until its preemption level is higher than the current system ceiling

Stack-Resource Policy


Rules of Basic Stack-Based, Preemption-Ceiling Protocol

0. *Update of the Current Ceiling:* Whenever all the resources are free, the preemption ceiling of the system is Ω . The preemption ceiling $\hat{\Psi}(t)$ is updated each time a resource is allocated or freed.
1. *Scheduling Rule:* After a job is released, it is blocked from starting execution until its preemption level is higher than the current ceiling $\Psi(t)$ of the system and the preemption level of the executing job. At any time t , jobs that are not blocked are scheduled on the processor in a priority-driven, preemptive manner according to their assigned priorities.
- ★ 2. *Allocation Rule:* Whenever a job J requests for a resource R , it is allocated the resource.
3. *Priority-Inheritance Rule:* When some job is blocked from starting, the blocking job inherits the highest priority of all the blocked jobs.

Here, “priority” stands for “deadline”, and preemption levels reflect task periods!!

- PCP: check system ceiling on locking
- SRP: check system ceiling on scheduling

Preemption
levels



Black=2
Gray=3




FIGURE 8-20 Example of priority ceilings of multiple unit-resources.

Stack-Resource Policy

- SRP avoids uncontrolled priority inversion
 - Protected by deadline-inheritance
- SRP avoids transitive blocking and chain blocking
 - When a job start executing, all that resources it needs have been released
 - Governed by system ceiling
- SRP prevent deadlocks from happening
 - A job is never be blocked once it start executing

Stack-Resource Policy

- Blocking time
 - Direct blocking (?)
 - No blocking on locking
 - Deadline-inheritance blocking
 - Blocked on scheduling
 - Preemption-ceiling blocking
 - Blocked on scheduling
- Use a similar technique for PCP/RM to calculate task locking times of SRP/EDF

Stack-Resource Policy

- A collection of tasks $\{T_1, T_2, \dots, T_n\}$ governed by SRP are schedulable by EDF if


$$\forall_{i=1}^n \left(\frac{b_i}{p_i} + U \right) \leq 1$$

- U = the total CPU utilization of all tasks
- The same test applies to NPCS and CPP


Stack-Resource Policy

- Pros:
 - Free from indirection priority inversion and deadlocks
 - Easy to implement, works with EDF
- Cons:
 - Response of HPTs is bad



Ungoverned



PCP+RM



SRP+EDF



| | Uncontrollable priority inversion | Bounded blocking time | Blocked at most once | Deadlock avoidance |
|---|--------------------------------------|--------------------------|-------------------------|-----------------------|
| Non-preemptible Critical Sections (NPCS) | No | Yes | Yes | Yes |
| Ceiling-Priority Protocol (CPP) | No | Yes | Yes | Yes |
| Priority- Inheritance Protocol (PIP) | No | Ouch | No | No |
| Priority-Ceiling Protocol (PCP) | No | Yes | Yes | Yes |
| Stack-Resource Policy (SRP) | No | Yes | Yes | Yes |

Summary

- Check list
 - Priority inversion
 - Blocking time
 - Blocking duration
 - Deadlock
 - Response

- Backup slices...

Priority-Ceiling Protocol

- Context-switch for independent tasks
 - Two context switches are accounted to the **preempting** job
 - Thus every job is added to $2*CS$
- With PCP, a job can be blocked for at most once
 - Additional $2*CS$ is accounted to a job with a **non-zero** blocking time

Priority-Ceiling Protocol

- A collection of tasks $\{T_1, T_2, \dots, T_n\}$ governed by PCP are schedulable if

$$\forall_{i=1 \dots n} \left\{ \frac{b_i}{p_i} + \sum_{j=1}^i \frac{f(j) \times CS + c_j}{p_j} \leq U(i) \right\}$$

- $f(x)=4$ if $b_x \neq 0$
- $f(x)=2$ if $b_x == 0$

Stack-Resource Policy


- Preemption ceilings of a **multi-instance** resources R increases as the free units in R decreases
- Let $\lceil R \rceil_k$ be the priority ceiling (preemption ceiling) of resource R there are k units available
 - $\lceil R \rceil_k$ is the highest preemption level of all tasks that require **more than k units ($>$)** of resource R
 - These task jobs may be blocked when trying to lock resource R

| Avail | | Need | | | | | | |
|-------|-------|------|----|----|---------------------|---------------------|---------------------|---------------------|
| R | N_R | J1 | J2 | J3 | $\lceil R \rceil_0$ | $\lceil R \rceil_1$ | $\lceil R \rceil_2$ | $\lceil R \rceil_3$ |
| R_1 | 3 | 3 | 2 | 1 | 3 | 2 | 1 | 0 |
| R_2 | 1 | 1 | 1 | 0 | 2 | 0 | 0 | 0 |
| R_3 | 3 | 1 | 3 | 1 | 3 | 2 | 2 | 0 |

Figure 4: Ceilings of Resources.

Job 1 need 3
units of R1

$\lceil R1 \rceil_2$: Ceiling of R1 when
there are two residual
instances of R1



Preemption levels:
J3 highest
J1 lowest

Stack-Resource Policy

- With PCP+RM, a task may be blocked (for once) in the middle of execution
 - Extra two units of the cxtsw overhead
- With SRP+EDF, a task will not be blocked once it starts execution
 - Last-In First-Out, this is why SRP is named after “stack” resource policy